

PURDUE UNIVERSITY

ECE 661 COMPUTER VISION

HOMEWORK 7

SUBMISSION: ARJUN KRAMADHATI GOPI

EMAIL: akramadh@purdue.edu

THEORY QUESTIONS

The reading material for Lecture 15 presents three different approaches to characterizing the texture in an image: 1) using the Grayscale Co- Occurrence Matrix (GLCM); 2) with Local Binary Pattern (LBP) histograms; and 3) using a Gabor Filter Family. Explain succinctly the core ideas in each of these three methods for measuring texture in images. (You are not expected to write more than a dozen sentences on each).

Grayscale Co-Occurrence Matrix - GLCM

- Raster scan image pixel by pixel - left to right and top to bottom.
- Set parameter d for the distance value.
- For each pixel in the scan, consider pixel at distance d .
- Increment the matrix GLCM at $[gl_1, gl_2]$ and $[gl_2, gl_1]$ by 1. Where gl_1 is the greylevel at the given pixel and gl_2 is the greylevel at the point at distance d .
- Characterise the matrix with factors like entropy, homogeneity and contrast. This will be the feature vector for the image used to train and classify.

Local Binary Patterns - LBP

- For each pixel in image, consider a circular boundary of radius R on which there are p number of points.
- Calculate greylevels at each of the p points on the circle.
- If greylevel at the point is lower than level at center then assign 0, else 1. Obtain rotation invariant pattern.
- Construct histogram of the frequency of the bin encoding. This will be the texture feature of the image.

Gabor Filter Family

- It banks on the property that most textures have repetitively occurring micro-patterns.
- Can be thought of as a highly localised Fourier transform in which the localization is achieved by applying a Gaussian decay function to the pixels.
- Gaussian weighting gives the localization of the pattern.
- The direction of the periodicities in the underlying Fourier kernel is used to characterize a texture in that direction.

*With regard to representing color in images, answer Right or Wrong for the following questions: (a) RGB and HSI are just linear variants of each other. (b) The color space $L^*a^*b^*$ is a nonlinear model of color perception. (c) Measuring the true color of the surface of an object is made difficult by the spectral composition of the illumination.*

- (a) RGB and HSI are just linear variants of each other - **RIGHT**
- (b) The color space $L^*a^*b^*$ is a nonlinear model of color perception - **RIGHT**
- (c) Measuring the true color of the surface of an object is made difficult by the spectral composition of the illumination - **RIGHT**

PROGRAMMING TASKS

The task is to build an image classification algorithm using the Local Binary Pattern (LBP) features and K-Nearest Neighbors (kNN) classifier. We will briefly discuss the following main concepts:

1. LBP algorithm
2. kNN algorithm

LOCAL BINARY PATTERNS (LBP)

The LBP algorithm is used to generate texture features from images. The texture features are invariant to transnational and rotational changes. To achieve this, we will have to accomplish the following tasks:

1. Extract binary patterns of inter-pixel variations in the image.
2. Generate rotation invariant representations from the Local Binary Patterns.
3. Encoding the minIntVal forms of the Local Binary Patters.

Extract binary patterns of the inter-pixel variations of the Image

The main idea behind this is to run 0s and 1s around a given pixel. What this means is that for each pixel, we define a circular boundary of radius **R** on which reside **P** number of points.




Figure 1: Image source: Dr. Avi Kak's tutorial on Modelling texture and color in images

In the figure above, we see 8 neighboring points on the circle of unit radius. We estimate the coordinates of these points using the following relation:

$$(\Delta u, \Delta v)_p = (R \cos \frac{2\pi p}{P}, R \sin \frac{2\pi p}{P}) \quad (1)$$

Where p ranges from 0 to $P-1$. Here, $P = 8$ and $R = 1$. We compute the coordinates of the points by adding the $\Delta u, \Delta v$ values to the x and y coordinates. Once we have the coordinates of these points on the circle, our job is to estimate the pixel values at those coordinates. But, it is evident that the estimated point coordinates will not exactly be an actual pixel coordinate in the image. Therefore, we will need to interpolate the pixel value at these coordinates using the 4 surrounding pixel values from actual pixels.




Figure 2: Image source: Dr. Avi Kak's tutorial on Modelling texture and color in images

We consider 4 pixels A, B, C and D as shown in the figure above. These pixels form a box which encapsulates the estimated point coordinate at p . Using the four known pixel values, we then interpolate the pixel value at p using the following relation:

$$greylevelatp = (1 - \Delta k)(1 - \Delta l)A + (1 - \Delta k)\Delta l B + \Delta k(1 - \Delta l)C + \Delta k\Delta l D \quad (2)$$

Using the above equation, we compute the grey levels at each of the 8 points on the circle. We then threshold these 8 values against the grey level at the center of the circle which is the pixel under consideration from the beginning. If the grey level at the point is greater than or equal to the level at the center then we assign a value of 1 to that point. If it is less than the grey level at the center we assign a value of 0 to the point. Thereby we obtain a binary vector for each of the pixel in the image. We use these binary vectors to generate rotation invariant representations in the next step.

Generate rotation invariant representations of the Local Binary Patterns

The idea behind generating a rotation invariant representation is to find the minIntVal by rotating the image around the pixel under consideration. When we rotate the image around the pixel at the center of the given circle, we will see that the points on the circle will also move in the same circular manner. These points simply shift from one place to another on the same circle. For each such shift (or rotation) we calculate the binary vector for the same center point/pixel. We then convert these binary representations to their decimal equivalent values. The least decimal value in this list would be that of the

binary pattern which is rotation invariant!

The idea behind this logic is derived from the idea that most number of 0's occupy the most significant positions when they are representing the rotation invariant form.

In this homework we make use of Dr. Avinash Kak's BitVector module on top of some own implementations to get the rotation invariant representations of the Local Binary Patterns.

Encoding the minIntVal forms of the Local Binary Patterns

In the previous step, we saw how we obtained the rotation invariant representation of the LBP at each pixel in the entire image. These representation are pixel wise representations. We need an encoding system to make this an image-wise representation from the pixel-wise representation that it is at the moment. According to the creators of the LBP, we can make use of the following rules to encode the pixels in the image based on their minIntValues. The rules are as follows:

1. If the binary sequence consists only of 0s we encode the pixel as 0
2. If the binary sequence consists only of 1s we encode the pixel as P. Where P is the total number of points on the circular vicinity.
3. If the number of runs is 2 then the encoding of the pixel is the number of 1s in the sequence.
4. If the number of runs is greater than 2 then the encoding of the pixel is $P+1$.

Using the encoding of every pixel in the entire image, we construct a normalised histogram which has $P+2$ bins starting from 0 and ending at the $P+1$ th bin.

This histogram will act as the feature vector for the image.

K-NEAREST NEIGHBORS (KNN)

Once we obtain the feature vectors of every image in the training data set, the question is how to use them to predict the image classes for the images in the testing data set.

For this, we make use of the kNN algorithm.

The idea is very simple. We store the feature vector of each image matched against the class of the image. For every new image we need predictions for:

1. We calculate the feature vector of the image.
2. We then estimate a list of k matches whose feature vectors resembles the feature vector of the image.
3. Out of the k matches, we chose the image class which repeats the most number of times in the list.
4. This label is then assigned to the given image and the prediction is deemed complete.

To compute the k nearest matches, we use a variety of distance calculations to measure the entries with the least distance. Euclidian distance, dot product distance and cosine distance are few of the popular ones. In this homework I have implemented the Euclidian based matching algorithm.

RESULTS




Figure 3: Class Beach - Sample LBP Histogram




Figure 4: Class Building - Sample LBP Histogram




Figure 5: Class Mountain - Sample LBP Histogram




Figure 6: Class Car - Sample LBP Histogram




Figure 7: Class Tree - Sample LBP Histogram

Confusion Matrix

	Beach	Building	Mountain	Car	Tree
Beach	4	0	0	1	0
Building	0	3	0	2	0
Mountain	0	1	4	0	0
Car	1	1	0	3	0
Tree	0	2	0	0	3

Accuracy calculation:

$$Accuracy = \frac{Truepredictions}{Totalpredictions} * 100 \quad (3)$$

We have True predictions = 17 from the confusion matrix Total Predictions have to be 25 since there are 5 images each in the five classes. Therefore, the accuracy is:

$$Accuracy = \frac{17}{25} * 100 = 68\% \quad (4)$$

Overall, we see that this is a decent image classifier. We used the parameters:

1. R value 1
2. P value 8
3. k value 5 for kNN Classifier

SOURCE CODE

Part of the kNN classifier was inspired from the implementation found in: [LINK](#)

```
1
2 """
3 Computer Vision - Purdue University - Homework 7
4 Author : Arjun Kramadhati Gopi, MS-Computer & Information
5         Technology, Purdue University.
6 Date: Oct 19, 2020
7
8 [TO RUN CODE]: python3 classifier.py
9 Output:
10 [labels]: Predictions for the input images
11 """
12
13 import cv2 as cv
14 from matplotlib import pyplot as plt
15 import numpy as np
16 import math
17 import BitVector
18 import pickle
19 import os
20 from collections import Counter
21
22
23 class Imageclassifier:
24
25     def __init__(self, training_directory, testing_directory,
26                 parameterR, parameterP, kvalue, Train=False):
27         """
28         Initialise the image classifier object used to either
29         train or test the classification of images
30         :param training_directory: Directory of the training set
31         :param testing_directory: Directory of the testing set
32         :param parameterR: Radius for the circular boundary
33         :param parameterP: Number of points on the circular
34                             boundary
35         :param kvalue:k nearest neighbors needed for match
36         :param Train: Training yes or no
37         """
38         self.classdict = dict()
39         self.imagedict = dict()
40         self.histogramdict = dict()
41         self.database = []
42         self.cmatrix = np.zeros((5, 5), dtype='int')
43         if not Train:
```



```

41         training_directory = testing_directory
42     self.classcount = len(os.listdir(training_directory))
43     for element in os.listdir(training_directory):
44         self.classdict[element] = len(os.listdir(
45             training_directory+'/'+element))
46     templist = []
47     for image in sorted(os.listdir(training_directory
48         ++'/'+element))):
49         print(os.listdir(training_directory+'/'+element))
50         print(image)
51         origimage = cv.imread(training_directory+'/'+
52             element+'/'+image)
53         imageread = np.zeros((origimage.shape[0],
54             origimage.shape[1], origimage.shape[2]), dtype
55             ='uint8')
56         image_gray = np.zeros((origimage.shape[0],
57             origimage.shape[1]), dtype='uint8')
58         imageread = cv.imread(training_directory+'/'+
59             element+'/'+image)
60         image_gray = cv.cvtColor(imageread, cv.
61             COLOR_BGR2GRAY)
62         templist.append(image_gray)
63     self.imagedict[element]=templist
64     self.parameterR = parameterR
65     self.parameterP = parameterP
66     self.kneighbors = kvalue
67     print(self.imagedict['beach'][0].shape)
68
69 def get_pixel_value(self, queuetuple, delu, delv, centerX,
70     centerY):
71     """
72     This function implements the bilinear interpolation
73     method used to get the pixel value
74     or the grey value at the point p
75     :param queuetuple: Location of the image
76     :param delu: change in value in x direction
77     :param delv: change in value in y direction
78     :param centerX: the point at the center of the circle
79         under consideration
80     :param centerY: the point at the center of the circle
81         under consideration
82     :return: greylevel at the point p
83     """
84     image = self.imagedict[queuetuple[0]][queuetuple[1]]
85     if (delu < 0.01) and (delv < 0.01):
86         interpolated_greylevel = float(image[centerX][centerY
87             ])
88     elif (delv < 0.01):
89         interpolated_greylevel = (1 - delu) * image[centerX][
90             centerY] + delu * image[centerX + 1][centerY]
91     elif (delu < 0.01):
92         interpolated_greylevel = (1 - delv) * image[centerX][
93             centerY] + delv * image[centerX][centerY + 1]

```



```

117         bitvector = BitVector.BitVector(bitlist=
118             binarypatternforpoint)
119         intvals_for_circular_shifts = [int(bitvector <<
120             1) for _ in range(self.parameterP)]
121         minimum_bit_vector = BitVector.BitVector(intVal=
122             min(intvals_for_circular_shifts), size=self.
123             parameterP)
124         runs = minimum_bit_vector.runs()
125         histogram = self.build_histogram(histogram, runs)
126     if Train:
127         self.histogramdict[queuetuple] = histogram
128         plt.bar(list(histogram.keys()), histogram.values(),
129             color='b')
130         path = 'histograms/' + str(queuetuple[0]) + '/'
131         plt.savefig(path + 'Class_{}'.format(queuetuple[0]) +
132             '_ImageNum_{}'.format(int(queuetuple[1])) + '.png'
133             ')
134     if not Train:
135         return histogram
136
137     def save_histograms_of_all(self, filename):
138         """
139         Saves the Local Binary Pattern histograms of every image
140         :param filename: File name for the data base
141         :return: None
142         """
143         file = open(filename, 'wb')
144         pickle.dump(self.histogramdict, file)
145         file.close()
146
147     def load_data(self, filename):
148         """
149         Loads the database from the saved .obj file. The database
150         contains
151         the LBP histograms of every image in training set
152         :param filename: File name of the database being
153             retrieved.
154         :return: Load and store the database in a dictionary
155         """
156         blist = []
157         bblist = []
158         clist = []
159         mlist = []
160         tlist = []
161         beachlist = np.zeros((20,10))
162         buildinglist = np.zeros((20,10))
163         carlist = np.zeros((20,10))
164         mountainlist = np.zeros((20,10))
165         treelist = np.zeros((20,10))
166         print(self.classdict['tree'])
167         file = open(filename, 'rb')
168         database = pickle.load(file)
169         file.close()

```

```

161         for element, index in database:
162             if element[0] == 'beach':
163                 blist.append(database.get(element))
164             if element[0] == 'building':
165                 bblast.append(database.get(element))
166             if element[0] == 'mountain':
167                 mlist.append(database.get(element))
168             if element[0] == 'car':
169                 clist.append(database.get(element))
170             if element[0] == 'tree':
171                 tlist.append(database.get(element))
172
173         for index in range(len(blist)):
174             beachlist[index, :] = np.array(list(blist[index].
175                 values()))
176         for index in range(len(bblast)):
177             buildinglist[index, :] = np.array(list(bblast[index].
178                 values()))
179         for index in range(len(clist)):
180             carlist[index, :] = np.array(list(clist[index].values
181                 ()))
182         for index in range(len(mlist)):
183             mountainlist[index, :] = np.array(list(mlist[index].
184                 values()))
185         for index in range(len(tlist)):
186             treelist[index, :] = np.array(list(tlist[index].
187                 values()))
188
189         histogram_all = np.zeros((100, 11))
190         for i in range(5):
191             index1 = 20 * i
192             index2 = index1 + 20
193             histogram_all[index1:index2, 0] = i
194             histogram_all[:, 1:] = np.concatenate((beachlist,
195                 buildinglist, carlist, mountainlist, treelist), axis
196                 =0)
197         self.database = histogram_all
198         print('loaded data successfully')
199
200     def knn_classify(self, list_histograms_class,
201         numberoftestingimages = 5, numberoftrainingimages =20):
202         """
203         This function implements the knnparameter-Nearest
204         Neighbor algorithm. We use the Euclidian distance to
205         calculate the nearest matches.
206         :param list_histograms_class: List of all the LBP
207             histograms of a particular class
208         :param numberoftestingimages: Number of images in the
209             testing set
210         :param numberoftrainingimages: Number of images in the
211             training set
212         :param nClass: Number of classes to predict
213         :return: returns the index of the label of the classes

```

```

202     """
203     knnparameter = self.kneighbors
204     training_histogram_all = self.database
205     result_hist = np.zeros((numberoftestingimages, 10))
206     condition1 = numberoftrainingimages * 1
207     condition2 = numberoftrainingimages * 2
208     condition3 = numberoftrainingimages * 3
209     condition4 = numberoftrainingimages * 4
210     condition5 = numberoftrainingimages * 5
211     for index_class in range(len(list_histograms_class)):
212         label_list = np.zeros((numberoftestingimages,
213                                knnparameter), dtype='int')
214         labelindex = np.zeros(numberoftestingimages, dtype='
215                                int')
216         result_hist[index_class, :] = np.array(list(
217             list_histograms_class[index_class].values()))
218         eucledian_distance = np.zeros((numberoftestingimages,
219                                         training_histogram_all.shape[0]))
220         for imageindex in range(numberoftestingimages):
221             for imagetrain in range(training_histogram_all.shape
222                                     [0]):
223                 eucledian_distance[imageindex, imagetrain] = np.
224                     linalg.norm(result_hist[imageindex, :] -
225                                 training_histogram_all[imagetrain, 1:])
226             sorted_distance = np.argsort(eucledian_distance[
227                 imageindex, :])
228             for k_idx in range(knnparameter):
229                 if (sorted_distance[k_idx] <= (condition1)):
230                     label_list[imageindex, k_idx] = 0
231                 elif (sorted_distance[k_idx] <= (condition2)):
232                     label_list[imageindex, k_idx] = 1
233                 elif (sorted_distance[k_idx] <= (condition3)):
234                     label_list[imageindex, k_idx] = 2
235                 elif (sorted_distance[k_idx] <= (condition4)):
236                     label_list[imageindex, k_idx] = 3
237                 elif (sorted_distance[k_idx] <= (condition5)):
238                     label_list[imageindex, k_idx] = 4
239             labelindex[imageindex], freq = Counter(list(
240                 label_list[imageindex, :])).most_common(1)[0]
241     return labelindex
242
243 def predict_and_analyse(self, blist, bblast, mlist, clist, tlist):
244     """
245     This function takes the histograms of the testing set and
246     uses them to
247     predict the class labels for each of the images in the
248     testing set. The results are
249     collated in a confusion matrix
250     :param blist: List of histograms for beach class
251     :param bblast: List of histograms for building class
252     :param mlist: List of histograms for mountain class
253     :param clist: List of histograms for the car class
254     :param tlist: List of histograms for the tree class

```

```

244         :return: Prints the final confusion matrix.
245         """
246
247         label_index = self.knn_classify(blist)
248         label_unique, label_unique_count = np.unique(label_index,
249             return_counts=True)
249         self.cmatrix[0, label_unique] = label_unique_count
250         label_index = self.knn_classify(bblist)
251         label_unique, label_unique_count = np.unique(label_index,
252             return_counts=True)
252         self.cmatrix[1, label_unique] = label_unique_count
253         label_index = self.knn_classify(mlist)
254         label_unique, label_unique_count = np.unique(label_index,
255             return_counts=True)
255         self.cmatrix[2, label_unique] = label_unique_count
256         label_index = self.knn_classify(clist)
257         label_unique, label_unique_count = np.unique(label_index,
258             return_counts=True)
258         self.cmatrix[3, label_unique] = label_unique_count
259         label_index = self.knn_classify(tlist)
260         label_unique, label_unique_count = np.unique(label_index,
261             return_counts=True)
261         self.cmatrix[4, label_unique] = label_unique_count
262         print('Prediction complete')
263         print('Printing confusion matrix...')
264         print(self.cmatrix)
265
266
267 if __name__ == "__main__":
268     """
269     Code begins here
270     """
271     tester = Imageclassifier("imagesDatabaseHW7/training", "
272         imagesDatabaseHW7/testing", 1, 8, 5)
272     Train = False
273     if Train:
274         for element in os.listdir("imagesDatabaseHW7/training"):
275             for index in range(len(os.listdir("imagesDatabaseHW7/
276                 training" + '/' + element))):
276                 print('training image class: '+element+' __ ' +
277                     str(index))
277                 tester.generate_texture_feature((element, index))
278             print('Training complete. Saving histogram dictionary
279                 ...')
279             tester.save_histograms_of_all('histograms.obj')
280             print('Saving complete')
281     tester.load_data('histograms.obj')
282     btestlist = []
283     bbtestlist = []
284     mtestlist = []
285     ctestlist = []
286     ttestlist = []
287     testdict = dict()

```

```
288     for element in os.listdir("imagesDatabaseHW7/testing"):
289         for index in range(len(os.listdir("imagesDatabaseHW7/
290             testing" + '/' + element))):
291             print('testing image class: ' + element + ' __ ' +
292                 str(index))
293             hist = tester.generate_texture_feature((element, index
294             ))
295             if element == 'beach':
296                 btestlist.append(hist)
297             if element == 'building':
298                 bbtestlist.append(hist)
299             if element == 'mountain':
300                 mtestlist.append(hist)
301             if element == 'car':
302                 ctestlist.append(hist)
303             if element == 'tree':
304                 ttestlist.append(hist)
305 tester.predict_and_analyse(btestlist, bbtestlist, mtestlist,
306                             ctestlist, ttestlist)
```